

2021F_Week05

September 22, 2021

0.1 Week05

Lets continue discussing the details of Regression. In particular, we can write down the maximum likelihood estimators for β_0 , β_1 , and σ^2 . Next we will write down the Fisher Information **Matrix** for a 2X1 vector that contains β_0 and β_1 .

With our MLE and Fisher Information Matrix we can compute confidence intervals for β and setup hypothesis tests.

Our goal this week is to: 1. Write down the solution to the loglikelihood 2. Write down and understand how to use the Fisher Information 3. Learn how to build Confidence intervals for β 4. Learn how to compute fitted values $\mu(\hat{x})$ 5. Briefly discuss the “true” model

0.1.1 The solution to the loglikelihood (ie the MLEs for β_0 , β_1 , and σ^2)

The MLE of β_1 ($\hat{\beta}_1$) is

$$\hat{\beta}_1 = \frac{\sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^N (x_i - \bar{x})^2} \quad (1)$$

The MLE of β_0 ($\hat{\beta}_0$) is

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} \quad (2)$$

The MLE for σ^2 is worth some discussion of the model form for SLR. Recall SLR model form is

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad (3)$$

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad (4)$$

In model form we see that given estimates for β_0 and β_1 we can compute ϵ_i values. Lets plug in our MLE estimates for β_0 and β_1 and solve for ϵ_i :

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad (5)$$

$$y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + \epsilon_i \quad (6)$$

$$\epsilon_i = y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i) \quad (7)$$

We can consider each $\epsilon_{-}\{i\}$ as a data point. Our data is $\epsilon = [\epsilon_1, \epsilon_2, \dots, \epsilon_N]$ and our model is

$$E_i \sim \mathcal{N}(0, \sigma^2) \quad (8)$$

We know from previous work that the MLE for σ^2 of this model is

$$\hat{\sigma}^2 = \frac{(\epsilon_i - \bar{\epsilon})^2}{N} \quad (9)$$

and by assumption, and the Law of Large Numbers, we expect $\bar{\epsilon} \approx 0$ so

$$\hat{\sigma}^2 = \frac{(\epsilon_i - 0)^2}{N} \quad (10)$$

$$\hat{\sigma}^2 = \frac{\epsilon_i^2}{N} \quad (11)$$

$$(12)$$

Though the above is the MLE for $\hat{\sigma}^2$ it is more traditional to estimate $\hat{\sigma}^2$ as

$$\hat{\sigma}^2 = \frac{\epsilon_i^2}{N-2} \quad (13)$$

$$(14)$$

0.1.2 Fitted values

Assume data $\mathcal{D} = [(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)]$ and the following model

$$Y_i | x_i \sim \mathcal{N}(\mu(x_i), \sigma^2) \quad (15)$$

where

$$\mu(x_i) = \beta_0 + \beta_1 x_i \quad (16)$$

Then $E(Y_i | x_i) = \mu(x_i) = \beta_0 + \beta_1 x_i$ and the MLE estimate of $\mu(x_i)$, denoted $\hat{\mu}(x_i)$ is called a **fitted value** for Y_i . A **fitted value** for Y_i assuming a SLR model is

$$\hat{\mu}(x_i) = \hat{\beta}_0 + \hat{\beta}_1 x_i \quad (17)$$

0.1.3 Treating parameters as a vector and the Fisher Information Matrix

Up until now we have specified our data and model for individual data points/random variables. Let us reformulate our data and model using Matrix algebra (because it is awesome).

Assume data $\mathcal{D} = [(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)]$ Our model says that each random variable Y_i has the following distribution

$$Y_1 = 1\beta_0 + x_1\beta_1 + \epsilon_1 \quad (18)$$

$$Y_2 = 1\beta_0 + x_2\beta_1 + \epsilon_2 \quad (19)$$

$$Y_3 = 1\beta_0 + x_3\beta_1 + \epsilon_3 \quad (20)$$

$$Y_4 = 1\beta_0 + x_4\beta_1 + \epsilon_4 \quad (21)$$

$$Y_5 = 1\beta_0 + x_5\beta_1 + \epsilon_5 \quad (22)$$

where each $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$.

This model can be rewritten as

$$Y = X\beta + \epsilon \quad (23)$$

where

$$Y = \begin{bmatrix} Y_1 \\ Y_2 \\ Y_3 \\ Y_4 \\ Y_5 \end{bmatrix} \quad (24)$$

$$\epsilon = \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \epsilon_4 \\ \epsilon_5 \end{bmatrix} \quad (25)$$

where each ϵ_i is normally distributed with parameter σ^2 or $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$.

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & x_3 \\ 1 & x_4 \\ 1 & x_5 \end{bmatrix} \quad (26)$$

and finally

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} \quad (27)$$

0.1.4 The Multivariate Normal distribution and probabilities over vectors

Because we have re-expressed our Simple linear regression using matrices and vectors, we need a way to talk about probabilities over vectors. To be clear, we rewrote our SLR

$$Y \sim \text{Some Distribution}(X\beta, \Sigma) \quad (28)$$

But how should we characterize a vector where each entry in the vector has a Normal Distribution?

The Multivariate Normal Distribution (MVN) The multivariate distribution uses a probability density over vectors to assign with assigning probabilities over intervals. For a **random vector** $x = [x_1, x_2, \dots, x_N]'$ the MVN Normal distribution is denoted

$$x \sim MVN(\mu, \Sigma) \quad (29)$$

where $\mu = [\mu_1, \mu_2, \dots, \mu_N]$ is a mean vector and

$$\Sigma = \begin{bmatrix} \sigma_{x_1}^2 & \sigma_{x_1, x_2} & \cdots & \sigma_{x_1, x_N} \\ \sigma_{x_1, x_2} & \sigma_{x_2}^2 & \cdots & \sigma_{x_2, x_N} \\ \vdots & \vdots & \ddots & \vdots \end{bmatrix} \quad (30)$$

is called the covariance matrix where diagonal entries in this matrix Σ_{ii} are equal to the variance of x_i and entries Σ_{ij} are equal to the covariance between x_i and x_j .

PDF The probability density for the multivariate normal distribution over vectors of length k is

$$f(x, \mu, \Sigma) = (2\pi)^{-k/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} [(x - \mu)' \Sigma^{-1} (x - \mu)] \right\} \quad (31)$$

The Expected value and Covariance For a random vector X of length k with a $MVN(\mu, \Sigma)$ distribution the expected value is

$$E(X) = \mu \quad (32)$$

where $\mu = [\mu_1, \mu_2, \dots, \mu_k]'$ is a vector of length k .

The covariance matrix (rather than variance) of a random variable X as above is

$$\text{Cov}(X) = \Sigma \quad (33)$$

For a vector of length two we can visualize the probability density function of the MVN for different parameter values.

```
[42]: import scipy.stats

x = np.linspace(-5, 5, 100)
y = np.linspace(-5, 5, 100)

X, Y = np.meshgrid(x, y)

def MVN(pos,m,s):
    return scipy.stats.multivariate_normal(mean=m,cov=s).pdf(pos)

pos = np.empty(X.shape + (2,))
pos[:, :, 0] = X
pos[:, :, 1] = Y

Z= MVN(pos,m=[0,0],s=[[2,0],[0,4]])

fig,axs = plt.subplots(1,3)

ax = axs[0]
ax.contour(X, Y, Z, colors='black');

ax.text(0.99,0.99,r"Cov=0",transform=ax.
    ↳transAxes,ha="right",va="top",fontsize=12)
ax.set_xlabel("X1")
ax.set_ylabel("X2")

Z= MVN(pos,m=[0,0],s=[[2,2],[2,4]])

ax = axs[1]
ax.contour(X, Y, Z, colors='black');

ax.text(0.99,0.99,r"Cov>0",transform=ax.
    ↳transAxes,ha="right",va="top",fontsize=12)
ax.set_xlabel("X1")
ax.set_ylabel("X2")

Z= MVN(pos,m=[0,0],s=[[2,-2],[-2,4]])
```


```

ax = axs[2]
ax.contour(X, Y, Z, colors='black');

ax.text(0.99,0.99,r"Cov<0",transform=ax.
    ↳transAxes,ha="right",va="top",fontsize=12)
ax.set_xlabel("X1")
ax.set_ylabel("X2")

fig.set_size_inches(18,4)

```



On the far left is a MVN density with 0 covariance between X_1 and X_2 , where

$$\begin{bmatrix} X_1 \\ X_2 \end{bmatrix} \sim MVN \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 2 & 0 \\ 0 & 4 \end{bmatrix} \right) \quad (34)$$

The middle figure is a MVN distribution with a positive covariance

$$\begin{bmatrix} X_1 \\ X_2 \end{bmatrix} \sim MVN \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 2 & 2 \\ 2 & 4 \end{bmatrix} \right) \quad (35)$$

and the figure on the far right is MVN with a negative covariance

$$\begin{bmatrix} X_1 \\ X_2 \end{bmatrix} \sim MVN \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 2 & -2 \\ -2 & 4 \end{bmatrix} \right) \quad (36)$$

0.1.5 Applying the MVN distributoin to our linear regression

Recall that our linear regression model can be written

$$Y = X\beta + \epsilon \quad (37)$$

where X is a matrix and Y , β , and ϵ are vectors.

Distribution for MLE of β The distribution over $\hat{\beta}$ is

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} \sim MVN \left(\hat{\beta}, \hat{\sigma}^2 (X'X)^{-1} \right) \quad (38)$$

where $\hat{\beta}$, the MLE for β_0 and for β_1 can be written using matrices

$$= (X'X)^{-1} X'y \quad (39)$$

Because $X\hat{\beta}$ is the μ we can write our fitted values \hat{y} as

$$\hat{y} = X \left[(X'X)^{-1} X'y \right] \quad (40)$$

$$= X \left[(X'X)^{-1} X' \right] y \quad (41)$$

$$= P_x y \quad (42)$$

$$(43)$$