

Week09

October 26, 2021

0.0.1 Multivariate regression

Up until now we have explored simple linear regression (SLR). SLR supposed a the conditional distribution of a random variable Y has a Normal distribution where the mean is a function of a random variable X and the variance is constant.

But often there may be many variables that change the probability distribution of Y .

Multivariate Linear Regression (MLR) is a statistical model that relates more than one random variable X (called covariates) to Y .

We will first rewrite SLR using a centered X random variable X^* to simplify maximum likelihood estimates of β_0, β_1 . Next, we will rewrite these MLEs using matrix algebra. Finally, we generalize these MLE estimating equations to more than one variable X .

Centering on X Our typical SLR setup assumes a dataset of x, y pairs generated from pairs of random variables and a specific probability distribution for Y . Suppose we collect a sample of N data points $\mathcal{D} = [(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)]$ which are generated from N pairs of random variables (X, Y) .

SLR supposes the following conditional probability distribution for Y

$$Y|x, \beta, \sigma^2 \sim \mathcal{N}(\beta_0 + \beta_1 x, \sigma^2) \quad (1)$$

or

$$Y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad (2)$$

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad (3)$$

Lets change the above SLR model by creating a new variable x^* by subtracting from each x_i the sample average \bar{x} . This is called centering on x and $x_i^* = x_i - \bar{x}$ is called a centered covariate.

The SLR model relating x^* and Y is then

$$Y_i = \beta_0 + \beta_1 x_i^* + \epsilon_i \quad (4)$$

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad (5)$$

We can visualize the effect of centering below

```

[60]: x      = np.random.normal(3,3,100) # generate 1000 points
      xbar   = np.mean(x)
      xstar  = x-xbar

      epsilons = np.random.normal(0,2,100)

      y= 2+3*x + epsilons

      fig,axs =plt.subplots(1,2)
      ax=axs[0]
      ax.scatter(x,y,alpha=0.7)
      ax.set(xlabel=r"$x$",ylabel="y")

      ax.axvline(0)
      ax.axhline(0)

      ax=axs[1]
      ax.scatter(xstar,y,alpha=0.7)
      ax.set(xlabel=r"$x^{\ast}$ (centered x)",ylabel="y")

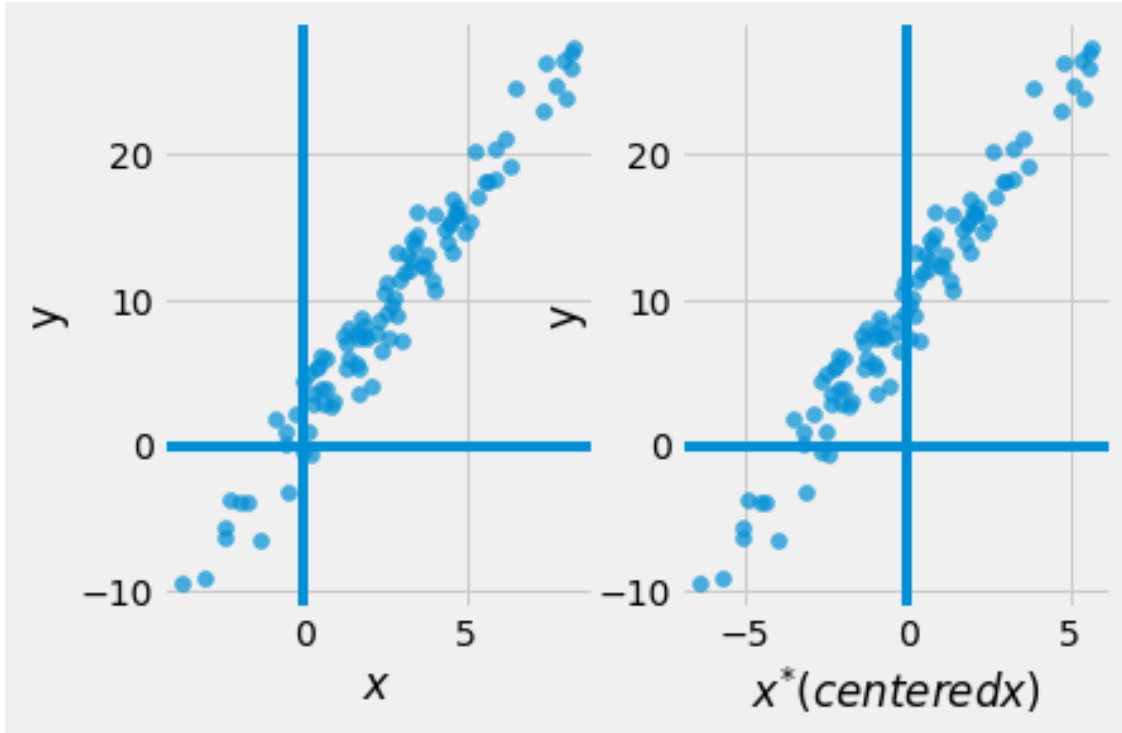
      ax.axvline(0)
      ax.axhline(0)

```

```

[60]: <matplotlib.lines.Line2D at 0x1c902af70>

```



By centering, all of the (x,y) pairs in this particular case move to the left.

In our original SLR, we can compute the MLEs for β_0 and β_1 as

$$\hat{\beta}_0 = \bar{y} - \beta_1 \bar{x} \quad (6)$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^N (x_i - \bar{x})^2} = \frac{\sum_{i=1}^N (x_i - \bar{x})y_i}{\sum_{i=1}^N (x_i - \bar{x})^2} \quad (7)$$

Our new SLR model where we center x simplifies the above estimates

$$\hat{\beta}_0^* = \bar{y} \quad (8)$$

$$\hat{\beta}_1^* = \frac{\sum_{i=1}^N x_i^* y_i}{\sum_{i=1}^N (x_i^*)^2} \quad (9)$$

Above we used the (very handy) fact that

$$\bar{x}^* = \frac{\sum_{i=1}^N (x_i - \bar{x})}{N} \quad (10)$$

$$= \frac{1}{N} \left(\sum_{i=1}^N x_i - \sum_{i=1}^N \bar{x} \right) \quad (11)$$

$$= \frac{1}{N} \left(\sum_{i=1}^N x_i - N\bar{x} \right) \quad (12)$$

$$= \frac{1}{N} (N\bar{x} - N\bar{x}) \quad (13)$$

$$= 0 \quad (14)$$

Are β_1 and β_1^* equivalent? Shifting all of our collected xvalues by the same constant does not change the relationship between Y and X. We can see that this is the case two ways.

Algebraic

$$Y_i = \beta_0^* + \beta_1^* x_i^* + \epsilon_i \quad (15)$$

$$= \beta_0^* + \beta_1^* (x_i - \bar{x}) + \epsilon_i \quad (16)$$

$$= \beta_0^* - \beta_1^* \bar{x} + \beta_1^* x_i + \epsilon_i \quad (17)$$

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad (18)$$

$$(19)$$

We see that our β_1^* is attached to x in the same way as our original SLR. This then means $\beta_1^* = \beta_1$

Visually

We can plot the points (x, y) and (x^*, y) and see that the estimated **slope** remains unchanged.

```
[61]: x      = np.random.normal(3,3,100) # generate 1000 points
xbar   = np.mean(x)
xstar  = x-xbar

epsilons = np.random.normal(0,2,100)

y= 2+3*x + epsilons

fig,axs =plt.subplots(1,2)
ax=axs[0]
ax.scatter(x,y,alpha=0.7)
ax.set(xlabel=r"$x$",ylabel="y")

minx,maxx = min(x),max(x)
b1,b0 = np.polyfit(x,y,1)
ax.plot([minx,maxx],[b0+b1*minx,b0+b1*maxx],color="red" )
```

```

ax.axvline(0)
ax.axhline(0)


ax=axis[1]
ax.scatter(xstar,y,alpha=0.7)
ax.set(xlabel=r"$x^{*}$ (centered x)",ylabel="y")

minx,maxx = min(xstar),max(xstar)
b1,b0 = np.polyfit(xstar,y,1)
ax.plot([minx,maxx],[b0+b1*minx,b0+b1*maxx],color="red" )

ax.axvline(0)
ax.axhline(0)

```

[61]: <matplotlib.lines.Line2D at 0x1c917fbe0>



Are β_0 and β_0^* equivalent? The intercept for the SLR and SLR with centered x are **different**.

From above, we saw that the SLR with centered x simplified to ##### Algebraic

$$Y_i = \beta_0^* - \beta_1^* \bar{x} + \beta_1^* x_i + \epsilon_i \quad (20)$$

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2) \quad (21)$$

$$(22)$$

The expression $\beta_0^* - \beta_1^* \bar{x}$ involve only constants. This expression is β_0 . That is, $\beta_0 = \beta_0^* - \beta_1^* \bar{x}$

Visually

We see from above that the estimated regression line (red line) intersects the y-axis at different points depending on whether we center or do not center our data.

Though centering makes computing MLEs easier this process also changes our β_0 or y-intercept.

0.0.2 Matrix algebra to compute the MLEs for centered SLR

Let us look again at the MLE estimates for sentered SLR

$$\hat{\beta}_0^* = \bar{y} \quad (23)$$

$$\hat{\beta}_1^* = \frac{\sum_{i=1}^N x_i^* y_i}{\sum_{i=1}^N (x_i^*)^2} \quad (24)$$

The MLE for $\hat{\beta}_1^*$ appears to be two inner products divided by one another. That is, we can compute $\hat{\beta}_1^*$ as

$$\hat{\beta}_1^* = \frac{x^* \cdot y}{x^* \cdot x^*} \quad (25)$$

where $x \cdot y = \sum_{i=1}^N x_i y_i$ and

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_N \end{bmatrix} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{bmatrix} \quad (26)$$

We can also recognize the MLE for $\hat{\beta}_0^*$ as an inner product

$$\hat{\beta}_0^* = \bar{y} \quad (27)$$

$$= \frac{1}{N} \sum_{i=1}^N y_i \quad (28)$$

$$= \frac{1}{N} \sum_{i=1}^N 1 \times y_i \quad (29)$$

$$= \frac{1}{N} \mathbf{1} \cdot y \quad (30)$$

where

$$\mathbf{1} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \quad (31)$$

Using matrix algebra, our MLEs become

$$\hat{\beta}_0^* = (\mathbf{1} \cdot \mathbf{1})^{-1}(\mathbf{1} \cdot y) \quad (32)$$

$$\hat{\beta}_1^* = (x^* \cdot x^*)^{-1}(x^* \cdot y) \quad (33)$$

A data setup such that we arrive at the above MLEs The above MLE estimates pop out of the following data setup. Let X be a matrix with N rows and 2 columns. The first column will be all ones and the second column will be filled, from top to bottom, with x_1, x_2, \dots, x_N .

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & x_3 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix} \quad (34)$$

The above MLE formulas may motivate us to try

$$(X'X)^{-1}X'y \quad (35)$$

$$X'y = \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ x_1 & x_2 & x_3 & \cdots & x_N \end{bmatrix} \times \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{bmatrix} = \begin{bmatrix} \mathbf{1}'y \\ x'y \end{bmatrix} \quad (36)$$

The first multiplication looks like a step in the right direction.

$$X'X = \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ x_1 & x_2 & x_3 & \cdots & x_N \end{bmatrix} \times \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & x_3 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix} = \begin{bmatrix} \mathbf{1}'\mathbf{1} & \mathbf{1}'x \\ \mathbf{1}'x & x'x \end{bmatrix} \quad (37)$$

If we centered our x covariate then \bar{x} is zero and so

$$\mathbf{1}'x = \sum_{i=1}^N x_i = N\bar{x} = 0 \quad (38)$$

We can use the above to simplify $X'X$

$$X'X = \begin{bmatrix} \mathbf{1}'\mathbf{1} & 0 \\ 0 & x'x \end{bmatrix} \quad (39)$$

The matrix $X'X$ is now a diagonal matrix—a matrix with non-zero elements for the entries (i, i) . The inverse of a diagonal matrix is easy to compute

$$(X'X)^{-1} = \begin{bmatrix} 1/\mathbf{1}'\mathbf{1} & 0 \\ 0 & 1/x'x \end{bmatrix} \quad (40)$$

and we can now see our final product is

$$(X'X)^{-1}X'y = \begin{bmatrix} 1/\mathbf{1}'\mathbf{1} & 0 \\ 0 & 1/x'x \end{bmatrix} \begin{bmatrix} \mathbf{1}'y \\ x'y \end{bmatrix} = \begin{bmatrix} \mathbf{1}'y/\mathbf{1}'\mathbf{1} \\ x'y/x'x \end{bmatrix} \quad (41)$$

We can see that the above matrix multiplication returns a vector with two MLEs: $\hat{\beta}_0^*$ and $\hat{\beta}_1^*$.

Variance for the MLEs follows a similar pattern. We saw in simple linear regression that the variance for $\hat{\beta}_1$ was

$$\mathbb{V}(\hat{\beta}_1) = \sigma^2 \frac{1}{\sum_{i=1}^N (x_i - \bar{x})^2} \quad (42)$$

and the variance for $\hat{\beta}_0$

$$\mathbb{V}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{N} + \frac{1}{\sum_{i=1}^N (x_i - \bar{x})^2} \right) \quad (43)$$

There is a natural generalization for the variance of $\hat{\beta}_0$ and $\hat{\beta}_1$ like the above generalization for the MLEs.

$$\mathbb{V}(\hat{\beta}) = \sigma^2 (X'X)^{-1} \quad (44)$$

The above variance will return a **covariance matrix** where the variance of $\hat{\beta}_0$ is the first diagonal entry and the variance of $\hat{\beta}_1$ is the second diagonal entry.

0.0.3 Multivariate Linear Regression (finally) and an example

Suppose we collect data points

$$\mathcal{D} = [(x_1^1, x_1^2, x_1^3, y_1), (x_2^1, x_2^2, x_2^3, y_2), \dots, (x_N^1, x_N^2, x_N^3, y_N)] \quad (45)$$

and propose that the conditional distribution of Y_i is

$$Y_i | x_i^1, x_i^2, \dots, x_i^3 \sim \mathcal{N}(\beta_0 + \beta_1 x_i^1 + \beta_2 x_i^2 + \beta_3 x_i^3, \sigma^2) \quad (46)$$

The above is **multivariate linear regression** model. We can compute the MLE for $\hat{\beta}$ as

$$\hat{\beta} = (X'X)^{-1}X'y \quad (47)$$

and we can compute the variance for our MLEs as

$$\mathbb{V}(\hat{\beta}) = \sigma^2(X'X)^{-1} \quad (48)$$

```
[62]: from sklearn.datasets import load_diabetes
import numpy.linalg as la

diabData = load_diabetes()
X = diabData["data"]
y = diabData["target"]

# add intercept (column of ones)
X = np.column_stack([np.ones((len(X),1)),X] )
```

```
[63]: betaHat = la.solve( X.T.dot(X), X.T.dot(y))
betaHat

epsilons = y - X.dot(betaHat)
sigma2Hat = epsilons.T.dot(epsilons) / len(epsilons)

varBetaHat = sigma2Hat * la.inv(X.T.dot(X))
```

```
[64]: import statsmodels.api as sm
model = sm.OLS(y,X)
res = model.fit().summary()
res
```

```
[64]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  y    R-squared:                  0.518
```

```

Model:                OLS      Adj. R-squared:      0.507
Method:              Least Squares    F-statistic:      46.27
Date:                Tue, 26 Oct 2021    Prob (F-statistic):    3.83e-62
Time:                00:27:29    Log-Likelihood:      -2386.0
No. Observations:    442    AIC:      4794.
Df Residuals:        431    BIC:      4839.
Df Model:            10
Covariance Type:      nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const         152.1335         2.576      59.061      0.000      147.071      157.196
x1            -10.0122        59.749     -0.168      0.867     -127.448      107.424
x2           -239.8191        61.222     -3.917      0.000     -360.151     -119.488
x3             519.8398        66.534       7.813      0.000      389.069      650.610
x4             324.3904        65.422       4.958      0.000      195.805      452.976
x5           -792.1842       416.684     -1.901      0.058     -1611.169       26.801
x6             476.7458       339.035       1.406      0.160     -189.621     1143.113
x7             101.0446       212.533       0.475      0.635     -316.685      518.774
x8             177.0642       161.476       1.097      0.273     -140.313      494.442
x9             751.2793       171.902       4.370      0.000      413.409     1089.150
x10            67.6254        65.984       1.025      0.306     -62.065      197.316
=====
Omnibus:                1.506    Durbin-Watson:           2.029
Prob(Omnibus):          0.471    Jarque-Bera (JB):         1.404
Skew:                   0.017    Prob(JB):                 0.496
Kurtosis:               2.726    Cond. No.                 227.
=====

```

Notes:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""

```

```

[65]: print(betaHat[1])
      print(np.sqrt(varBetaHat[1,1]))

```

```

-10.012197817471035
59.00101633853383

```

```

[ ]:

```